

What is...?

Julia is an open-source, multi-platform, high-level, high-performance programming language for technical computing.

Julia has an LLVM-based JIT compiler that allows it to match the performance of languages such as C and FORTRAN without the hassle of low-level code. Because the code is compiled on the fly you can run (bits of) code in a shell or REPL, which is part of the recommended workflow.

Julia is dynamically typed, provides multiple dispatch , and is designed for parallelism and distributed computation.

Julia has a built-in package manager.

Julia has many built-in mathematical functions, including special functions (e.g. Gamma), and supports complex numbers right out of the box.

Julia allows you to generate code automatically thanks to Lisp-inspired macros.

Julia was born in 2012.

Basics

Assignment	answer = 42 x, y, z = 1, [1;10],], "A string" x, y = y, x # swap x and y
Constant declaration	const DATE_OF_BIRTH = 2012
End-of-line comment	! = 1 # This is a comment
Delimited comment	#= This is another comment =#
Chaining	x = y = z = 1 # right-to-left 0 < x < 3 # true 5 < x != y < 5 # false
Function definition	function add_one(i) return i + 1 end
Insert LaTeX symbols	\delta + [Tab]

Operators

Basic arithmetic	+, -, *, /
Exponentiation	2^3 == 8
Division	3/12 == 0.25
Inverse division	7\3 == 3/7
Remainder	x % y or rem(x,y)
Negation	!true == false
Equality	a == b
Inequality	a != b or a ≠ b
Less and larger than	< and >
Less than or equal to	<= or ≤
Greater than or equal to	>= or ≥
Element-wise operation	[1, 2, 3] .+ [1, 2, 3] == [2, 4, 6] [1, 2, 3] .* [1, 2, 3] == [1, 4, 9]
Not a number	isnan(NaN) !not(!) NaN == NaN
Ternary operator	a == b ? "Equal" : "Not equal"
Short-circuited AND and OR	a && b and a b
Object equivalence	a == b

The shell a.k.a. REPL

Recall last result	ans
Interrupt execution	[Ctrl] + [C]
Clear screen	[Ctrl] + [L]
Run program	include("filename.jl")
Get help for func is defined	?func
See all places where func is defined	apropos("func")
Command line mode	:
Package Manager mode]
Help mode	?
Exit special mode / Return to REPL	[Backspace] on empty line
Exit REPL	exit() or [Ctrl] + [D]

Standard libraries

To help Julia load faster, many core functionalities exist in standard libraries that come bundled with Julia. To make their functions available, use using PackageName. Here are some Standard Libraries and popular functions.

Random	rand, randn, randnsubseq
Statistics	mean, std, cor, median, quantile
LinearAlgebra	I, eigvals, eigvecs, det, cholseq
SparseArrays	sparse, SparseVector, SparseMatrixCSC
Distributed	@distributed, pmap, addprocs
Dates	DateTime, Date

Package management

Packages must be registered before they are visible to the package manager. In Julia 1.0, there are two ways to work with the package manager: either with using Pkg and using Pkg functions, or by typing] in the REPL to enter the special interactive package management mode. (To return to regular REPL, just hit BACKSPACE on an empty line in package management mode). Note that new tools arrive in interactive mode first, then usually also become available in regular Julia sessions through Pkg module.

Using Pkg in Julia session

List installed packages (human-readable)	Pkg.status()
List installed packages (machine-readable)	Pkg.installed()
Update all packages	Pkg.update()
Install PackageName	Pkg.add("PackageName")
Rebuild PackageName	Pkg.build("PackageName")
Use PackageName (after install)	using PackageName
Remove PackageName	Pkg.rm("PackageName")

In Interactive Package Mode

Add PackageName	add PackageName
Remove PackageName	rm PackageName
Update PackageName	update PackageName
Use development version	dev PackageName or dev gitrepoUrl
Stop using development version, revert to public release	free PackageName

Characters and strings

Character	chr = 'C'
String	str = "A string"
Character code	Int('J') == 74
Character from code	Char(74) == 'J'
Any UTF character	chr = '\uXXXX' # 4-digit HEX chr = '\UXXXXXXXX' # 8-digit HEX
Loop through characters	for c in str println(c) end
Concatenation	str = "Learn" * " " * "Julia"
String interpolation	a = b = 2 println("a * b = \$(a*b)", "Julia")
First matching character or regular expression	findFirst(isequal('i'), "Julia")
Replace substring or regular expression	replace("Julia", "a" => "us") == "Julius"
Last index (of collection)	lastindex("Hello") == 5
Number of characters	length("Hello") == 5
Regular expression	pattern = r"l[aeiou]" str = "1 234 567 890" pat = r"\s*([0-9]) ([0-9]+)" m = match(pat, str) m.captures == ["1", "234"] [m.match for m = eachmatch(pat, str)]
All occurrences	eachmatch(pat, str)
All occurrences (as iterator)	eachmatch(pat, str)
Beware of multi-byte Unicode encodings in UTF-8:	10 == lastindex("Ångström") != length("Ångström") == 8
Strings are immutable.	

Numbers

Integer types	IntN and UIntN, with N ∈ {8, 16, 32, 64, 128}, BigInt
Floating-point types	FloatN with N ∈ {16, 32, 64}, BigFloat
Minimum and maximum values by type	typemin(Int8), typemax(Int64)
Complex types	Complex{T}
Imaginary unit	im
Machine precision	eps{C} # same as eps(Float64)
Rounding	round{C} # floating-point round(Int, x) # Integer
Type conversions	convert{TypeName, val} # attempt/error typename(val) # calls convert
Global constants	pi # 3.1415... n # 3.1415... im # real(im * im) == -1
More constants	using Base.MathConstants
Julia does not automatically check for numerical overflow. Use package SafeIntegers for ints with overflow checking.	

Random Numbers

Many random number functions require using Random.

Set seed	seed!(seed)
Random numbers	rand{C} # uniform [0,1) randn{C} # normal (-Inf, Inf)
Random from Other Distribution	using Distributions my_dist = Bernoulli(0.2) # For example rand(my_dist)
Random subsample elements from A with inclusion probability p	randsubseq(A, p)
Random permutation elements of A	shuffle(A)

Arrays

Declaration	arr = Float64[]
Pre-allocation	sizehint!(arr, 10^4)
Access and assignment	arr = Any{1,2} arr[1] = "some text" a = [1;10;] b = a a[1] = 99 # b points to a a == b a == b # true
Copy elements (not address)	b = copy(a) b = deepcopy(a)
Select subarray from m to n	arr[m:n]
n-element array with 0.0s	zeros(n)
n-element array with 1.0s	ones(n)
n-element array with #undefs	Vector{Type}(undef,n)
n equally spaced numbers from start to stop	range(start, stop=stop, length=n)
Array with n random Int8 elements	rand{Int8, n}
Fill array with val	fill!(arr, val)
Pop last element	pop!(arr)
Pop first element	popfirst!(a)
Push val as last element	push!(arr, val)
Push val as first element	pushfirst!(arr, val)
Remove element at index idx	deleteat!(arr, idx)
Sort	sort!(arr)
Append a with b	append!(a,b)
Check whether val is element	in(val, arr) or val in arr
Scalar product	dot(a, b) == sum(a .* b)
Change dimensions (if possible)	reshape{1:6, 3, 2}' == [1 2 3; 4 5 6]
To string (with delimiter del between elements)	join(arr, del)

Linear Algebra

For most linear algebra tools, use using LinearAlgebra.

Identity matrix	I # just use variable I. Will automatically conform to dimensions required.
Define matrix	M = [1 0; 0 1]
Matrix dimensions	size(M)
Select i th row	M[i, :]
Select i th column	M[:, i]
Concatenate horizontally	M = [a b] or M = hcat(a, b)
Concatenate vertically	M = [a ; b] or M = vcat(a, b)
Matrix transposition	transpose(M)
Conjugate matrix transposition	M' or adjoint(M)
Matrix trace	tr(M)
Matrix determinant	det(M)
Matrix rank	rank(M)
Matrix eigenvalues	eigvals(M)
Matrix eigenvectors	eigvecs(M)
Matrix inverse	inv(M)
Solve M*x = v	M\v is better than inv(M)*v
Moore-Penrose pseudo-inverse	pinv(M)
Julia has built-in support for matrix decompositions.	
Julia tries to infer whether matrices are of a special type (symmetric, hermitian, etc.), but sometimes fails. To aid Julia in dispatching the optimal algorithms, special matrices can be declared to have a structure with functions like Symmetric, Hermitian, UpperTriangular, LowerTriangular, Diagonal, and more.	

Control flow and loops

Conditional	if-elseif-else-end
Simple for loop	for i in 1:10 println(i) end
Unnested for loop	for i in 1:10, j = 1:5 println(i+j) end
Enumeration	for (idx, val) in enumerate(arr) println("the \$idx-th element is \$val") end
while loop	while bool_expr # do stuff end
Exit loop	break
Exit iteration	continue

Functions

All arguments to functions are passed by reference.

Functions with ! appended change at least one argument, typically the first: sort!(arr).

Required arguments are separated with a comma and use the positional notation.

Optional arguments need a default value in the signature, defined with =.

Keyword arguments use the named notation and are listed in the function's signature after the semicolon:

```
function func(req1, req2; key1=df1t1, key2=df1t2)  
    # do stuff  
end
```

The semicolon is *not* required in the call to a function that accepts keyword arguments.

The return statement is optional but highly recommended.

Multiple data structures can be returned as a tuple in a single return statement.

Command line arguments julia script.jl arg1 arg2... can be processed from global constant ARGS:

```
for arg in ARGS  
    println(arg)  
end
```

Anonymous functions can best be used in collection functions or list comprehensions: x -> x^2.

Functions can accept a variable number of arguments:

```
function func(a...)  
    println(a)  
end
```

```
func(1, 2, [3;5]) # tuple: (1, 2, UnitRange{Int64}{3;5})
```

Functions can be nested:

```
function outerfunction()  
    # do some outer stuff  
    function innerfunction()  
        # do inner stuff  
        # can access prior outer definitions  
    end  
    # do more outer stuff  
end
```

Functions can have explicit return types

```
# take any number subtype and return it as a String  
function stringifynumber(num::T)::String where T <: Number  
    return "Snm"
```

Dictionaries

Dictionary	d = Dict{key1 => val1, key2 => val2, ...} d = Dict{key1 => val1, :key2 => val2, ...}
All keys (iterator)	keys(d)
All values (iterator)	values(d)
Loop through key-value pairs	for (k,v) in d println("key: \$k, value: \$v") end
Check for key :k	haskey(d, :k)
Copy keys (or values) to array	arr = [k for (k,v) in d]
Dictionaries are mutable; when symbols are used as keys, the keys are immutable.	

Sets

Declaration	s = Set{1, 2, 3, "Some text"}]
Union s1 U s2	union(s1, s2)
Intersection s1 ∩ s2	intersect(s1, s2)
Difference s1 \ s2	setdiff(s1, s2)
Difference s1 Δ s2	symdiff(s1, s2)
Subset s1 ⊆ s2	issubset(s1, s2)
Checking whether an element is contained in a set is done in O(1).	

Collection functions

Apply f to all elements of collection coll	map(f, coll) or map(coll) do elem # do stuff with elem end
Filter coll for true values of f	filter(f, coll)
List comprehension	arr = [f(elem) for elem in coll]

Types

Julia has no classes and thus no class-specific methods.

Types are like classes without methods.

Abstract types can be subtyped but not instantiated.

Concrete types cannot be subtyped.

By default, struct s are immutable.

Immutable types enhance performance and are thread safe, as they can be shared among threads without the need for synchronization.

Objects that may be one of a set of types are called Union types.

Type annotation	var::TypeName struct Programmer name::String birth_year::UInt16 fave_language::AbstractString end
Type declaration	
Mutable type declaration	mutable struct
Type alias	const Nerd = Programmer
Type constructors	methods{TypeName}
Type instantiation	me = Programmer("Ian", 1984, "Julia") me = Nerd("Ian", 1984, "Julia")
Subtype declaration	abstract type Bird end struct Duck <: Bird pond::String end
	struct Point{T :: Real} x::T y::T end
Parametric type	
	p = Point{Float64}(1,2)
Union types	Union{Int, String}
Traverse type hierarchy	supertype{TypeName} and subtypes{TypeName}
Default supertype	Any
All fields	fieldnames{TypeName}
All field types	TypeName.types

When a type is defined with an *inner* constructor, the default *outer* constructors are not available and have to be defined manually if need be. An inner constructor is best used to check whether the parameters conform to certain (invariance) conditions. Obviously, these invariants can be violated by accessing and modifying the fields directly, unless the type is defined as immutable. The new keyword may be used to create an object of the same type.

Type parameters are invariant, which means that Point{Float64} <: Point{Real} is false, even though Float64 <: Real. Tuple types, on the other hand, are covariant: Tuple{Float64} <: Tuple{Real}.

The type-inferred form of Julia's internal representation can be found with code_type{T}(). This is useful to identify where Any rather than type-specific native code is generated.

Missing and Nothing

Programmers Null	nothing
Missing Data	missing
Not a Number in Float	NaN
Filter missings	collect(skipmissing([1, 2, missing])) == [1,2]
Replace missings	collect{(:df[:col], 1)}
Check if missing	ismissing(x) not x == missing

Exceptions

Throw SomeExcept	throw(SomeExcept())
Rethrow current exception	rethrow()
Define NewExcep	Base.showerror(io::IO, e::NewExcep) = print(io, "A problem with \$(e.v):") throw(NewExcep("x"))
Throw error with msg text	error(msg)
	try # do something potentially iffy catch ex if isa(ex, SomeExcept) # handle SomeExcept elseif isa(ex, AnotherExcep) # handle AnotherExcep else # handle all others end finally # do this in any case end

Modules

Modules are separate global variable workspaces that group together similar functionality.

Definition	module PackageName # add module definitions # use export to make definitions accessible end
Include filename.jl	include("filename.jl")
Load	using ModuleName # all exported names using ModuleName: x, y # only x, y import ModuleName # only ModuleName import ModuleName: x, y # only x, y import ModuleName.x, ModuleName.y # only x, y # Get an array of names exported by Module names{ModuleName}
Exports	# include non-exports, deprecateds # and compiler-generated names names{ModuleName, all::Bool}
	#also show names explicitly imported from other modules names{ModuleName, all::Bool, imported::Bool}

With using Foo you need to say function Foo.bar(...) to extend module Foo's function bar with a new method, but with import Foo.bar, you only need to say function bar(...) and it automatically extends module Foo's function bar.

Expressions

Julia is homoiconic: programs are represented as data structures of the language itself. In fact, everything is an expression Expr.

Symbols are interned strings prefixed with a colon. Symbols are more efficient and they are typically used as identifiers, keys (in dictionaries), or columns in data frames. Symbols cannot be concatenated.

Quoting :(...) or quote ... end creates an expression, just like Meta.parse(str), and Expr{::call, ...}.

```
x = 1  
line = "1 + 5x" # some code  
expr = Meta.parse(line) # make an Expr object  
typeof(expr) == Expr # true  
dump(expr) # generate abstract syntax tree  
eval(expr) == 2 # evaluate Expr object: true
```

Macros

Macros allow generated code (i.e. expressions) to be included in a program.

Definition	macro macroname(expr) # do stuff end
Usage	macroname(ex1, ex2, ...) or @macroname ex1, ex2, ...
	@test # equal (exact) @ttest x = y # isapprox(x, y) @assert # assert (unit test) @which # types used @time # time and memory statistics @elapsed # time elapsed @allocated # memory allocated @profile # profile @spawn # run at some worker @spawnat # run at specified worker @asynct # asynchronous task @distributed # parallel for loop @everywhere # make available to workers

Rules for creating *hygienic* macros:

- Declare variables inside macro with local.
- Do not call eval inside macro.
- Escape interpolated expressions to avoid expansion: \$(esc(expr))

Parallel Computing

Parallel computing tools are available in the Distributed standard library.

Launch REPL with N workers	julia -p N
Number of available workers	nprocs()
Add N workers	addprocs(N)
See all worker ids	for pid in workers() println(pid) end
Get id of executing worker	myid()
Remove worker	rmprocs(pid) r = remotecall(f, pid, args...) # or: r = @spawnat pid f(args) ... fetch(r)
Run f with arguments args on pid (more efficient)	remotecall_fetch(f, pid, args...)
Run f with arguments args on any worker	r = @spawn f(args) ... fetch(r)
Run f with arguments args on all workers	r = [@spawnat w f(args) for w in workers()] ... fetch(r)
Make expr available to all workers	@everywhere expr
Parallel for loop with reducer function red	sum = @distributed{red} for i in 1:10 # do stuff end
Apply f to all elements in collection coll	pmap(f, coll)

Workers are also known as concurrent/parallel processes.

Modules with parallel processing capabilities are best split into a functions file that contains all the functions and variables needed by all workers, and a driver file that handles the processing of data. The driver file obviously has to import the functions file.

A non-trivial (word count) example of a reducer function is provided by Adam DeConinck.

I/O

Read stream	stream = stdin for line in eachline(stream) # do stuff end
Read file	open(filename) do file for line in eachline(file) # do stuff end end
Read CSV file	using CSV data = CSV.read(filename)
Write CSV file	using CSV CSV.write(filename, data)
Save Julia Object	using JLD save(filename, "object_key", object, ...)
Load Julia Object	using JLD d = load(filename) # Returns a dict of objects
Save HDF5	using HDF5 h5write(filename, "key", object)
Load HDF5	using HDF5 h5read(filename, "key")

DataFrames

For dplyr-like tools, see DataFramesMeta.jl.

Read Stata, SPSS, etc.	StatFiles Package
Describe data frame	describe(df)
Make vector of column col	v = df[:col]
Sort by col	sort!(df, [:col])
Categorical cols	categorical!(df, [:col])
List col levels	levels(df[:col])
All observations with col==val	df[df[:col] .== val, :]
Reshape from wide to long format	stack(df, [:lin, :] melt(df, [:col1, :col2], ...)
Reshape from long to wide format	unstack(df, :id, :val)
Make Nullable	allowmissing!(df) or allowmissing{df, :col} for r in eachrow(df) # r is struct with fields of col names. end
Loop over Rows	for c in eachcol(df) # do stuff. end
Loop over Columns	vector end by(df, :group_col, func) using Query query = @from r.col1 > df @where r.col1 > 40 @select (new_name=:col1, r.col2) @collect DataFrame # Default: iterator

Introspection and reflection

Type	typeof(name)
Type check	isa(name, TypeName)
List subtypes	subtypes{TypeName}
List supertype	supertype{TypeName}
Function methods	methods{Func}
JIT bytecode	code_llvm{expr}
Assembly code	code_native{expr}

Noteworthy packages and projects

Many core packages are managed by communities with names of the form Julia[Topic].

Statistics	JuliaStats
Differential Equations	JuliaDiffEq (DifferentialEquations.jl)
Automatic Differentiation	JuliaDiff
Numerical optimization	JuliaOpt
Plotting	JuliaPlots
Network (Graph) Analysis	JuliaGraphs
Web	JuliaWeb
Geo-Spatial	JuliaGeo
Machine Learning	JuliaML
	DataFrames # linear/logistic regression Distributions # Statistical distributions Flux # Machine learning Gadfly # ggplot2-likeplotting LightGraphs # Network analysts TextAnalysis # NLP

Naming Conventions

The main convention in Julia is to avoid underscores unless they are required for legibility.

Variable names are in lower (or snake) case: somevariable.

Constants are in upper (or snake) case: SOMECONSTANT.

Functions are in lower (or snake) case: somefunction.

Macros are in lower (or snake) case: @somemacro.

Type names are in initial-capital camel case: SomeType.

Julia files have the jl extension.

For more information on Julia code style visit the manual: [style guide](#).

Performance tips

- Avoid global variables.
- Write type-stable code.
- Use immutable types where possible.
- Use sizehint! for large arrays.
- Free up memory for large arrays with arr = nothing.
- Access arrays along columns, because multi-dimensional arrays are stored in column-major order.
- De-allocate resultant data structures.
- Disable the garbage collector in real-time applications: disable_gc{C}.
- Avoid the splat (...) operator for keyword arguments.
- Use mutating APIs (i.e. functions with ! to avoid copying data structures.
- Use array (element-wise) operations instead of list comprehensions.
- Avoid try-catch in (computation-intensive) loops.
- Avoid Any in collections.
- Avoid abstract types in collections.
- Avoid string interpolation in I/O.
- Vectorizing does not improve speed (unlike R, MATLAB or Python).
- Avoid eval at run-time.

IDEs, Editors and Plug-ins

- Julia (editor)
- JuliaBox (online Julia notebook)
- Jupyter (online Julia notebook)
- Emacs Julia (editor)
- vim Julia mode (editor)
- VS Code extension (editor)

Resources

- Official documentation.
- Learning Julia page.
- Month of Julia
- Community standards.
- Julia: A Fresh approach to numerical computing (pdf)
- Julia: A Fast Dynamic Language for Technical Computing (pdf)

Videos

- The 5th annual JuliaCon 2018
- The 4th annual JuliaCon 2017 (Berkeley)
- The 3rd annual JuliaCon 2016
- Getting Started with Julia by Leah Hanson
- Intro to Julia by Huda Nassar
- Introduction to Julia for Pythonists by John Pearson